

Response to Reviewers

Manuscript: R4Tun: LLM-guided adaptive segmental tunnel lining segmentation in point clouds

Manuscript number: AUTCON-D-26-02407

We are grateful to the Associate Editor and reviewers for their time in reviewing the manuscript and for their useful and insightful feedback. We have responded to the queries individually below and modified the manuscript as appropriate.

Summary of major revisions. The revised manuscript:

- reframes R4Tun as a *mechanism-contribution* study (bounded, cross-LLM, zero-label parameter adaptation) rather than a deployment-ready segmentation method;
- adds an error-analysis subsection (Section 4.6; Figs. 11–12; Tables 9–10) decomposing false negatives, false positives, and class swaps, and quantifying four structural limits of the fixed SAM4Tun labelling substrate;
- adds external positioning against supervised and expert-tuned tunnel-segmentation methods (Section 5.2; Table 11);
- reports a runtime trade-off against the base SAM4Tun processing time (Section 4.4; Table 7);
- adds a repeatability appendix (Appendix I; Table 25);
- expands limitations in the Conclusion, distinct from future-work directions.

All changes are highlighted in the revised manuscript PDF.

Associate Editor

Comment:

Avoid orphan headings without context: these are headers that are immediately followed by the next header, without any text in between (line 85–86, 169–170, ...). At least add a brief introduction to explain what will be covered in the section.

Response:

Short introductory paragraphs have been added at the opening of Sections 2–5 (Related work, Methodology, Results, and Discussion), each summarising what the section covers.

Comment:

Fix reference on line 191. Also in other places, the reference to the Appendix is missing. All references in the PDF need to resolve to a number. The Appendix can be directly included in the manuscript after the references.

Response:

Unresolved appendix references have been corrected. All appendix cross-references now resolve to numbered appendices (A–I). The Appendix is included directly in the manuscript after the References.

Comment:

Limitations must be included and discussed in the conclusion. Limitations are also not necessarily the same thing as Future Work, as some limitations are simply inherently part of the proposed approach (and thus need to be highlighted in a conclusion as well).

Response:

Section 6 (Conclusions) now contains a dedicated limitations paragraph, separate from the future-work directions that follow. The limitations discussed include:

- evaluation on SAM4Tun and Seg2Tunnel only;
- single expert-tuned reference anchor;
- inherited SAM4Tun operator assumptions and failure modes;
- dependence on offline bounds and knowledge quality;
- single-pass adaptation without iterative self-correction.

Comment:

Do a thorough check of all references such that they follow the author guidelines and include full bibliographic details. Author names, journal or book titles, chapter or article titles, year of publication, volume numbers, article numbers or pagination must be included, where applicable. Several names are mistyped, likely due to improper handling in L^AT_EX. Include DOIs for articles and ISBNs for book volumes. Check whether arXiv preprint articles have not been published by now, and try to replace them in that case, such that you can rely more on peer-reviewed and validated sources rather than pre-prints. Be careful with capitalization in titles as well, when using L^AT_EX; many acronyms are now in lower-case. When citing an article with 6+ authors, only the first six authors need to be listed, and the remaining ones can be replaced by et al. More guidance on this can be found in the author guide of the journal, and the examples it refers to.

Response:

The reference list has been checked for author names, article titles, journal titles, volume/article numbers, page ranges, DOIs, capitalisation, and published versions of preprints where available. L^AT_EX capitalisation errors for acronyms (e.g., SAM, LLM, GPR) have been corrected, long author lists truncated to six authors plus “et al.” where required, and DOIs added for journal articles.

Reviewer 1**Comment:**

The manuscript is generally well revised and significantly improved. One minor suggestion is to further strengthen the discussion of recent studies on intelligent tunnel inspection, adaptive parameter tuning, and AI-assisted point-cloud analysis for underground infrastructure. In particular, the authors may consider citing several recent related works to better position the contribution of R4Tun within the latest development of LLM-/AI-enabled tunnel analysis and infrastructure inspection research:

- DOI: [10.1016/j.aei.2026.104774](https://doi.org/10.1016/j.aei.2026.104774)
- DOI: [10.1016/j.autcon.2026.106769](https://doi.org/10.1016/j.autcon.2026.106769)
- DOI: [10.1016/j.tust.2024.105998](https://doi.org/10.1016/j.tust.2024.105998)
- DOI: [10.1016/j.measurement.2023.113903](https://doi.org/10.1016/j.measurement.2023.113903)

Response:

We thank the reviewer for this positive assessment. We have added the four suggested references to Related work and Discussion:

- Yue et al. (AEI 2026; DOI 10.1016/j.aei.104774) on non-contact infrastructure surface inspection;
- Yue et al. (AUTCON 2026; DOI 10.1016/j.autcon.106769) on AI-assisted damage segmentation in heritage buildings;
- Yue et al. (TUST 2024; DOI 10.1016/j.tust.105998) on shield-tunnel lining damage mechanisms;
- Yue et al. (Measurement 2024; DOI 10.1016/j.measurement.113903) on GPR-based tunnel-lining defect recognition.

These citations better position R4Tun within recent AI-enabled tunnel and infrastructure inspection research.

Comment:

The proposed framework currently evaluates adaptation quality mainly through segmentation metrics (mIoU and OA). It would be helpful if the authors could briefly discuss whether the adapted parameters remain stable across repeated runs or different API calls, especially considering that LLM-based reasoning may still introduce slight stochasticity even under low-temperature settings. A short discussion on reproducibility would further strengthen the practical applicability of the framework.

Response:

We added a repeatability analysis in Appendix I (LLM inference repeatability; Table 25) and a brief discussion in Section 4.3. Under $m + s + k$ with temperature set to 0, we compared run 1 and run 2 parameter JSONs on held-out complex-tunnel pairs. On the first nine completed tunnel-LLM pairs (tunnels 4-4, 5-3, and 5-4, three LLMs each):

- median critical-parameter identity across the 18 shared parameters was 100% (mean 88%);
- mean $|\Delta mIoU|$ between the two runs was 0.029 (median 0.000).

The full protocol targets 90 pairs (30 tunnels $\times$ 3 LLMs); aggregate statistics will be updated when the remaining pairs complete. Large mIoU swings on isolated pairs coincide with lower parameter identity and reflect failure-recovery reruns rather than benign stochasticity at temperature 0.

Comment:

The paper would benefit from a short clarification regarding the computational overhead introduced by the multi-agent adaptation process. Although runtime and API costs are reported, a brief comparison between the adaptation time and the original SAM4Tun inference time would help readers better understand the trade-off between adaptability and efficiency in practical deployment scenarios.

Response:

We have added a direct runtime comparison in Section 4.4 (Runtime trade-off; Table 7). We separated the base SAM4Tun processing time (235 s per tunnel) from the additional LLM adaptation time under $m + s + k$:

LLM	Extra time	Total time	Mean $\Delta mIoU$
Gemini-3-Flash	+96 s (0.4 $\times$)	$\sim$ 331 s	+0.166
Opus-4.6	+140 s (0.6 $\times$)	$\sim$ 375 s	+0.192
GPT-5.4	+307 s (1.3 $\times$)	$\sim$ 542 s	+0.171

Comment:

In Section 5.2, the limitations are discussed clearly. As a minor improvement, the authors could also briefly comment on how the proposed framework might generalize to other point-cloud processing pipelines beyond SAM4Tun, for example whether the context-design strategy (memory-state-knowledge) is expected to remain transferable even when the downstream segmentation backbone changes.

Response:

Section 5.3 (Limitations) acknowledges that transferability to other pipelines is conceptualised without empirical validation. The $m + s + k$ design can apply to other parameter-sensitive multi-stage point-cloud pipelines if four conditions hold:

- bounded tunable stage parameters;
- reference configurations with comparable target characterisation;
- summarisable intermediate state;
- encodable stage-specific knowledge.

The downstream segmentation backbone may change provided these conditions are met for the new pipeline stages.

Reviewer 2

Comment:

I would like to commend the authors for their efforts in revising this manuscript. The authors have clearly taken the previous reviewers' comments. The expansion of the dataset to 30 tunnel subsets, the introduction of a deterministic non-LLM rule-based baseline, the rigorous cumulative ablation studies, and the inclusion of comprehensive statistical analyses have significantly elevated the methodological rigor of the paper. The concept of utilizing LLMs as white-box reasoning agents to adapt expert-designed pipelines via structured context is interesting.

Response:

We thank the reviewer for this positive assessment and for highlighting the concept of our approach. We have clarified the focus of our study within the Abstract, Introduction, and Conclusion of the revised manuscript.

Comment:

However, despite the exhaustive experimental design, I have a fundamental and critical concern regarding the performance of the proposed system. The empirical results, as currently presented, fall drastically short of the standards required for practical deployment in construction automation, which severely undermines the core premise of the paper.

Response:

We agree that the absolute metrics are below deployment thresholds for autonomous inspection. We have therefore reframed the study as an exploration of a controlled, cross-LLM, zero-label, no-retraining, no-expert-tuned self-adaptation *mechanism*, rather than a performance-competitive segmentation method. Operating in a single forward pass (~ 331 s to ~ 542 s per tunnel), R4Tun introduces explicit, auditable LLM reasoning into an expert-designed pipeline and serves as a methodological foundation for future development.

Comment:

Massive Misclassification: *An OA of $\sim 53\%$ implies that nearly half of the point cloud is incorrectly classified. An mIoU of $\sim 33\%$ indicates a drastic lack of overlap between the predicted segment masks and the ground truth.*

Response:

We agree that these absolute figures are low. Section 4.6 (Error analysis) now decomposes where the errors arise. On regular tunnels, adaptation largely removes false negatives ($21\% \rightarrow 2\%$ of ground-truth points) but class swaps persist ($26\% \rightarrow 21\%$); on complex tunnels, adaptation recovers under-segmented blocks ($66\% \rightarrow 17\%$ false negatives) but recovered points are often mislabelled (class swaps rise to 34%). False positives remain small ($\leq 6\%$) throughout. The residual error is therefore dominated by the fixed positional-labelling rule (K-anchor placement and segment-offset template), not by LLM parameter selection.

Comment:

Failure in Downstream Tasks: *In the context of automated tunnel inspection (e.g., structural health monitoring, deformation analysis, water leakage detection, or Scan-to-BIM workflows), a segmentation pipeline with a $\sim 50\%$ error rate is practically unusable. The downstream geometric and structural analyses would be entirely compromised by the massive amount of noise, missing segments, and misidentified structural components.*

Response:

We agree. The revised manuscript states explicitly—in the Introduction, Discussion (Section 5.1), and Conclusion—that the current absolute accuracy does not yet support autonomous

final inspection or downstream tasks such as structural health monitoring, deformation analysis, water-leakage detection, or Scan-to-BIM without human verification.

Comment:

***The Low Baseline Fallacy:** The authors celebrate a $>100\%$ relative improvement, but this is a classic statistical illusion caused by a catastrophically failing baseline. The static SAM4Tun baseline completely collapses on off-reference tunnels (overall $mIoU = 0.15$; complex tunnels = 0.04). Recovering from a state of total failure to a state of mediocre/poor performance ($mIoU\ 0.34$) does not validate the framework for real-world engineering applications.*

Response:

We agree that headline relative-gain framing can be misleading when the baseline is near zero. We have removed celebratory relative-improvement language (e.g., “ $>100\%$ relative improvement”) from the Abstract, Discussion, and Conclusion. Discussion now leads with absolute $mIoU$ and OA; paired absolute deltas (e.g., $0.15 \rightarrow 0.32\text{--}0.34$) are retained only where they aid interpretation of the adaptation effect. We do not present recovery from baseline collapse as evidence of deployment readiness.

Comment:

***Poor Performance on Regular Tunnels:** Most concerningly, even on the regular tunnel subsets—which closely resemble the expert-tuned reference configuration—the $mIoU$ only reaches $\sim 0.50\text{--}0.54$. If an LLM, armed with explicit memory, intermediate state, and expert knowledge documents, cannot achieve a high degree of accuracy (e.g., $mIoU > 0.80$) on in-distribution or near-reference data, it raises serious doubts about either the LLM’s actual reasoning capability in this domain or the fundamental brittleness of the underlying SAM4Tun pipeline.*

Response:

We agree this is a fair concern. Section 4.6 shows that on regular tunnels the LLM recovers detection (false negatives drop from 21% to 2%) but cannot eliminate class swaps ($26\% \rightarrow 21\%$), which trace to the fixed K-anchor and segment-offset template rather than incorrect parameter choices. The near-reference $mIoU$ of $0.50\text{--}0.54$ therefore reflects a *substrate ceiling* of the single-reference SAM4Tun labelling rule, not a failure of LLM reasoning over structured context. This is consistent with the Non-LLM vs. LLM comparison: on regular tunnels, deterministic rules do not beat the static baseline, whereas the LLM ($m + s + k$) does.

Comment:

***Q1. Justification of Practical Applicability:** The primary scope of this journal is automation in construction. How do the authors justify the practical deployment of a system that misclassifies roughly 50% of the structural point cloud? The authors must provide a realistic discussion on what specific, limited engineering tasks (if any) an $mIoU$ of 0.34 can actually support, rather than broadly claiming it “reduces manual per-tunnel re-tuning.”*

Response:

We agree that broad deployment claims are not supported at $mIoU \approx 0.34$. We have removed claims of “reduces manual per-tunnel re-tuning” throughout the revised manuscript and state explicitly that current accuracy does not support autonomous final inspection.

Within this scope, we discuss limited tasks that the framework *may* support today, subject to human review:

- first-pass, bounded parameter adaptation before expert verification, producing a JSON parameter record and logged rationale;
- flagging when a target tunnel departs far from the calibrated reference anchor;
- shifting expert effort from repeated per-tunnel intervention toward upfront calibration and knowledge authoring (Discussion, Section 5.1).

We do *not* claim fitness for structural health monitoring, deformation analysis, water-leakage detection, or Scan-to-BIM without human verification.

Comment:

Q2. In-depth Error Analysis: *The manuscript currently lacks a qualitative and quantitative error analysis. Where are these massive errors concentrated? Are the LLMs failing to identify the correct boundaries? Is the 2D-to-3D reprojection introducing severe artifacts? Are background noises being classified as key blocks? The authors must provide visual error maps (e.g., False Positives vs. False Negatives) and analyze why the absolute metrics remain so stubbornly low despite LLM intervention.*

Response:

We agree. A new error-analysis subsection, Section 4.6 (Error analysis), has been added. It reports:

- FP/FN error maps on representative regular (tunnel 1-1) and complex (tunnel 4-1) examples (Fig. 11);
- a quantitative error-composition table over all ground-truth points (Table 9);
- quantified diagnostics for four structural limits of the fixed SAM4Tun labelling substrate (Table 10; Fig. 12): ring-level non-uniform point density, a moving K-anchor, fixed segment-offset template drift, and hard-coded walk direction.

The Discussion (Section 5.1) links these failure modes to the stated absolute performance ceiling and to the mechanism-level interpretation of the adaptation results. False positives (background classified as blocks) remain rare ($\leq 6\%$); the dominant residual errors are class swaps from the positional labelling rule, not 2D-to-3D reprojection artefacts or widespread noise misclassification as key blocks.

Comment:

Q3. Comparison with State-of-the-Art (SOTA) Baselines: *The current comparison is strictly internal (LLM vs. Non-LLM vs. Static SAM4Tun). How does an absolute mIoU of 0.34 compare to existing, traditional geometric feature-engineering methods or standard supervised 3D deep learning models trained on similar datasets? If a standard, non-LLM baseline can achieve an mIoU of 0.65+, the justification for introducing the computational overhead and API costs of LLMs is severely weakened.*

Response:

We agree that internal comparisons alone are insufficient. We added Section 5.2 (Positioning relative to supervised and expert-tuned methods) and Table 11, which situates R4Tun against established paradigms on Seg2Tunnel:

- supervised 3D deep learning (PointNet++, RandLA-Net, voxel-based methods): mIoU $\approx$ 0.83–0.93, but requires per-point labels and GPU retraining;
- expert-tuned SAM4Tun: mIoU >0.90 on curated test rings, but requires per-case expert tuning;
- geometric feature engineering: cross-section fitting error ~ 1 mm RMSE reported, but per-segment mIoU not directly comparable;
- single-reference static SAM4Tun (this study): mIoU 0.15 without re-tuning.

R4Tun (mIoU 0.32–0.34) is well below these methods on absolute accuracy, and we do not claim otherwise. The comparison is on setup cost: R4Tun is the only entry in Table 11 that is simultaneously label-free, training-free, and expert-tuning-free while still adapting per tunnel. We therefore frame the paper as a mechanism study, not a ready-to-deploy replacement for supervised or expert-tuned segmentation.

Comment:

***Q4. The Ceiling of the Pipeline:** The authors briefly mention the “single-reference design” as a limitation. However, if the underlying SAM4Tun pipeline is inherently incapable of achieving high mIoU due to its 3D-to-2D projection and SAM-prompting mechanics, then using an LLM to tune its parameters is merely “rearranging deck chairs on the Titanic.” The authors need to clarify whether the low mIoU is a failure of the LLM’s parameter selection or a fundamental limitation of the SAM4Tun architecture itself.*

Response:

We agree this distinction must be stated clearly. Section 4.6 identifies four structural constraints within the fixed SAM4Tun positional-labelling substrate:

1. **Non-uniform point density:** central rings are far denser than end rings (up to $19\times$ regular, $38\times$ complex); a single tunnel-level setup cannot span both.
2. **Moving K-anchor:** staggered assembly shifts K along the arc; K-mislocated rings rise from 12% (regular) to 64% (complex), collapsing per-ring accuracy.
3. **Fixed segment-offset template:** each block is placed at a fixed angular offset from K; the 6-segment template fits regular tunnels but fails on 7-segment complex geometry.
4. **Hard-coded walk direction:** reversed-handedness rings produce mirror-image labels that parameter tuning cannot correct.

The low absolute mIoU is therefore primarily a limitation of this labelling substrate, not incorrect LLM parameter selection. Bounded adaptation recovers detection (false-negative reduction) but cannot fix mirror flips or 7-segment template mismatch without rewriting the labelling stage—which lies outside the scope of bounded-parameter adaptation studied here. These constraints define the configuration challenge that future dynamic, context-aware labelling must address; they do not invalidate the mechanism-level finding that structured context enables cross-LLM parameter adaptation within fixed bounds.

Comment:

The integration of LLMs for context-aware parameter adaptation is interesting. The authors have done a good job structuring the experiments and ablation studies. Unfortunately, the actual numerical outcomes ($mIoU \sim 0.33$, $OA \sim 0.53$) reveal that the resulting system is riddled with segmentation errors and is currently unfit for practical, real-world tunnel inspection. Unless the authors can provide a compelling, evidence-based justification for why these low absolute metrics are acceptable, or demonstrate that the errors are benign for specific downstream tasks, the manuscript in its current form does not meet the practical impact threshold expected by the readership of Automation in Construction.

Response:

We agree that the system is not fit for autonomous real-world tunnel inspection at current accuracy. To address this concern, we have revised the manuscript to frame the contribution honestly and to supply the evidence requested above. Summary of actions:

- **Reframed study scope:** repositioned as a mechanism-contribution study; R4Tun explores zero-label, no-retraining, no-expert-tuned self-adaptation and does not yet support autonomous final inspection.
- **Performance framing:** removed celebratory relative-gain language; discussion leads with absolute mIoU and OA and acknowledges the performance ceiling.
- **In-depth error analysis:** Section 4.6 (Figs. 11–12; Tables 9–10) decomposes errors and quantifies four structural bottlenecks of the fixed labelling substrate.
- **External benchmarking:** Section 5.2 and Table 11 situate R4Tun against supervised and expert-tuned tunnel-segmentation methods.
- **LLM vs. rule distinction:** Section 3.4 (agent design, context, CoT), Algorithm 1, and Section 3.5.2 (Non-LLM baseline) separate LLM reasoning from deterministic rule lookup.
- **Limitations and future work:** Conclusion lists inherent limitations separately from future directions (dynamic labelling, multiple reference anchors, iterative self-correction).

Reviewer 3

Comment:

The study objectives and general rationale are mostly clear. However, the role and generality of the LLM in the proposed framework require more precise framing. The current wording may suggest that R4Tun provides a general LLM-based tunnel segmentation capability, whereas the implemented workflow still relies on an expert-designed pipeline, including tunnel unfolding, denoising, image enhancement, SAM-based segmentation, and 3D reprojection. Please revise the motivation, especially around Fig. 1, to clarify that the main contribution is LLM-guided bounded parameter adaptation for an existing segmentation pipeline, rather than a generally capable LLM solution for tunnel segmentation.

Response:

We agree. Fig. 1 and its caption have been revised to clarify that R4Tun extends a fixed expert-designed pipeline (SAM4Tun) via bounded parameter adaptation, rather than providing a general LLM-based tunnel segmentation capability. The Abstract, Introduction, and Conclusion use the same framing throughout.

Comment:

The method is generally described in sufficient detail, but some clarification would improve reader understanding. Additional explanation is needed to clarify how the LLM converts tunnel characteristics, reference parameters, intermediate pipeline states, and stage-specific knowledge documents into meaningful bounded parameter updates. Without this explanation, the specific role of LLM-based reasoning remains difficult to distinguish from deterministic rule-based parameter adjustment.

Response:

We agree. Section 3.4 (Agent design, Context design, CoT design) and Algorithm 1 describe how each stage agent assembles memory, state, and knowledge into a prompt and runs a five-step CoT protocol. Section 3.5.2 and Appendix E define the deterministic Non-LLM baseline (same knowledge encoded as explicit if-then lookups on characterisation fields).

Section 4.1, Table 4, and Fig. 9 compare Non-LLM vs. LLM: on regular tunnels, rules do not beat the static baseline; LLM ($m + s + k$) does; $m + s$ exceeds the rule baseline by 0.10–0.12 mIoU overall.

Appendix C gives the denoising agent’s memory, state, and knowledge inputs; Appendix D gives a full five-step CoT trace and returned JSON for one tunnel.

Comment:

The manuscript would benefit from improving several existing figures and tables so that they are more self-contained.

- **Fig. 1:** The motivation in Fig. 1 should be revised to avoid overstating the role of the LLM. The figure should clarify that R4Tun reduces repeated deployment-stage expert retuning by reusing expert-defined reference parameters, parameter bounds, and knowledge documents, rather than removing expert knowledge from the workflow.
- **Fig. 2:** Some labels, such as B1 and K2, are difficult to understand from the figure alone. These labels should be defined directly in the figure or caption so that readers do not need to refer back to the main text.
- **Fig. 5:** The relationship between the three interpolation methods shown on the left and the workflow shown on the right is ambiguous. The figure should clarify whether these methods are used sequentially or presented as alternatives.
- **Table 7:** Some parameters, especially the underlined parameters, are not sufficiently explained. The authors should add short descriptions of what these parameters control and how they affect the corresponding stages of the segmentation pipeline.

Response:

Fig. 1: revised to clarify that R4Tun reuses expert-defined reference parameters, parameter bounds, and knowledge documents, emphasising reduced deployment-stage retuning within a predefined pipeline rather than removal of expert knowledge.

Fig. 2: revised figure and caption define B1, K, and related labels directly.

Fig. 5: revised figure and caption clarify that the interpolation methods are sequential steps in one workflow.

Table 7: short descriptions added for the underlined parameters, explaining what each controls and how it affects the corresponding pipeline stage.

Comment:

The strengths of the study are present, but they could be emphasised with greater precision. The manuscript should more clearly state that the strength of R4Tun lies in integrating feature-engineered preprocessing, a foundation segmentation model, and LLM-guided bounded parameter adaptation. This framing would better highlight the actual contribution and avoid overstating the LLM as a general-purpose tunnel segmentation model.

Response:

We have revised the Abstract, Introduction (contributions list), and Conclusion to state explicitly that R4Tun integrates feature-engineered preprocessing, a foundation segmentation model (SAM4Tun), and LLM-guided bounded parameter adaptation—not a standalone LLM segmentation capability.

Comment:

The limitations are generally stated, but the manuscript should more explicitly acknowledge that R4Tun does not eliminate the need for expert knowledge. Expert input is still required to define reference parameters, parameter bounds, and stage-specific knowledge documents.

Response:

We have made this explicit in Section 5.3 (Limitations) and Section 6 (Conclusions). The revised manuscript presents the LLM as adapting within expert-defined constraints (reference parameters, bounds, and knowledge documents), not replacing them.

Comment:

The manuscript structure is generally clear, but the framing and figure explanations require revision to improve flow.

Response:

We have revised the Abstract, Introduction, and Conclusion to reframe the scope of the framework. We have also revised the captions of Figs. 1, 2, and 5 to improve readability and manuscript flow.

We believe the revised manuscript now presents R4Tun honestly as a controlled adaptation mechanism, reports absolute performance without overstated relative-gain framing, and provides the error decomposition and external positioning requested by Reviewer 2. We welcome any further guidance from the Editor and reviewers.